

Agenda

- File IO
- Stream API

Java IO framework

- Input/Output functionality in Java is provided under package java.io and java.nio package.
- IO framework is used for File IO, Network IO, Memory IO, and more.
- Two types of APIs are available file handling
 - FileSystem API -- Accessing/Manipulating Metadata
 - File IO API -- Accessing/Manipulating Contents/Data

Java IO

- Java File IO is done with Java IO streams.
- Java IO Streams are completely different from java.util.Stream. No relation between them
- Stream generally determines flow of data
- Java supports two types of IO streams.
 - Byte streams (binary files) -- byte by byte read/write
 - Character streams (text files) -- char by char read/write
- Stream is abstraction of data source/sink.
 - Data source -- InputStream(Byte Stream) or Reader(Char Stream)
 - Data sink -- OutputStream(Byte Stream) or Writer(Char Stream)
- All these streams are AutoCloseable (so can be used with try-with-resource construct)

Chaining IO Streams

- Each IO stream object performs a specific task.
 - FileOutputStream -- Write the given bytes into the file (on disk).
 - BufferedOutputStream -- Hold multiple elements in a temporary buffer before flushing it to underlying stream/device. Improves performance.
 - DataOutputStream -- Convert primitive types into sequence of bytes. Inherited from DataOutput interface.
 - ObjectOutputStream -- Convert object into sequence of bytes. Inherited from ObjectOutput interface.
 - PrintStream -- Convert given input into formatted output.
 - Note that input streams does the counterpart of OutputStream class hierarchy.
- Streams can be chained to fulfil application requirements.

Primitive types IO

- DataInputStream & DataOutputStream -- convert primitive types from/to bytes
 - primitive type --> DataOutputStream --> bytes --> FileOutputStream --> file.
 - DataOutput interface provides methods for conversion - writeInt(), writeUTF(), writeDouble(), ...
 - primitive type <-- DataInputStream <-- bytes <-- FileInputStream <-- file.

- DataInput interface provides methods for conversion - readInt(), readUTF(), readDouble(), ...

DataOutput/DataInput interface

- interface DataOutput
 - writeUTF(String s)
 - writeInt(int i)
 - writeDouble(double d)
 - writeShort(short s)
 - ...
- interface DataInput
 - String readUTF()
 - int readInt()
 - double readDouble()
 - short readShort()
 - ...

Serialization

- ObjectOutputStream & ObjectOutputStream -- convert java object from/to bytes
 - Java object --> ObjectOutputStream --> bytes --> FileOutputStream --> file.
 - ObjectOutputStream interface provides method for conversion - writeObject().
 - Java object <-- ObjectInputStream <-- bytes <-- FileInputStream <-- file.
 - ObjectInput interface provides methods for conversion - readObject().
- Converting state of object into a sequence of bytes is referred as Serialization. The sequence of bytes includes object data as well as metadata.
- Serialized data can be further saved into a file (using FileOutputStream) or sent over the network (Marshalling process).
- Converting (serialized) bytes back to the Java object is referred as Deserialization.
- These bytes may be received from the file (using FileInputStream) or from the network (Unmarshalling process).

ObjectOutput/ObjectInput interface

- interface ObjectOutput extends DataOutput
 - writeObject(obj)
- interface ObjectInput extends DataInput
 - obj = readObject()

Serializable interface

- Object can be serialized only if class is inherited from Serializable interface; otherwise writeObject() throws NotSerializableException.
- Serializable is a marker interface.

File

- File is a collection of data and information on a storage device.
- File = Data + Metadata
- collection of data/info on storage disk
- data = contents
- metadata = Information

java.io.File class

- A path (of file or directory) in file system is represented by "File" object.
- Used to access/manipulate metadata of the file/directory.
- Provides FileSystem APIs
 - `String[] list()` -- return contents of the directory
 - `File[] listFiles()` -- return contents of the directory
 - `boolean exists()` -- check if given path exists
 - `boolean mkdir()` -- create directory
 - `boolean mkdirs()` -- create directories (child + parents)
 - `boolean createNewFile()` -- create empty file
 - `boolean delete()` -- delete file/directory
 - `boolean renameTo(File dest)` -- rename file/directory
 - `String getAbsolutePath()` -- returns full path (drive:/folder/folder/...)
 - `String getPath()` -- return path
 - `File getParentFile()` -- returns parent directory of the file
 - `String getParent()` -- returns parent directory path of the file
 - `String getName()` -- return name of the file/directory
 - `static File[] listRoots()` -- returns all drives in the systems.
 - `long getTotalSpace()` -- returns total space of current drive
 - `long getFreeSpace()` -- returns free space of current drive
 - `long getUsableSpace()` -- returns usable space of current drive
 - `boolean isDirectory()` -- return true if it is a directory
 - `boolean isFile()` -- return true if it is a file
 - `boolean isHidden()` -- return true if the file is hidden
 - `boolean canExecute()`
 - `boolean canRead()`
 - `boolean canWrite()`
 - `boolean setExecutable(boolean executable)` -- make the file executable
 - `boolean setReadable(boolean readable)` -- make the file readable
 - `boolean setWritable(boolean writable)` -- make the file writable
 - `long length()` -- return size of the file in bytes
 - `long lastModified()` -- last modified time
 - `boolean setLastModified(long time)` -- change last modified time

transient fields

- `writeObject()` serialize all non-static fields of the class. If fields are objects, then they are also serialized.
- If any field is intended not to serialize, then it should be marked as "transient".

- The transient and static fields (except serialVersionUID) are not serialized.

serialVersionUID field

- Each serializable class is associated with a version number, called a serialVersionUID.
- It is recommended that programmer should define it as a static final long field (with any access specifier). Any change in class fields expected to modify this serialVersionUID.

```
private static final long serialVersionUID = 1001L;
```

- During deserialization, this number is verified by the runtime to check if right version of the class is loaded in the JVM. If this number mismatched, then InvalidClassException will be thrown.
- If a serializable class does not explicitly declare a serialVersionUID, then the runtime will calculate a default serialVersionUID value for that class (based on various aspects of the class described in the Java(TM) Object Serialization specification).

Buffered streams

- Each write() operation on FileOutputStream will cause data to be written on disk (by OS). Accessing disk frequently will reduce overall application performance. Similar performance problems may occur during network data transfer.
- BufferedOutputStream classes hold data into a in-memory buffer before transferring it to the underlying stream. This will result in better performance.
 - Java object --> ObjectOutputStream --> BufferedOutputStream --> FileOutputStream --> file on disk.
- Data is sent to underlying stream when buffer is full or flush() called explicitly.
- BufferedInputStream provides a buffering while reading the file.
- The buffer size can be provided while creating the respective objects.

PrintStream class

- Produce formatted output (in bytes) and send to underlying stream.
- Formatted output is done using methods print(), println(), and printf().
- System.out and System.err are objects of PrintStream class.
- It is used only to write the formatted data in to the file.

Scanner class

- Added in Java 5 to get the formatted input.
- It is java.util package (not part of java io framework).

```
Scanner sc = new Scanner(inputStream);  
// OR  
Scanner sc = new Scanner(inputFile);
```

- Helpful to read text files line by line.

Character streams

- Character streams are used to interact with text file.
- Java char takes 2 bytes (unicode), however char stored in disk file may take 1 or more bytes depending on char encoding.
 - <https://www.w3.org/International/questions/qa-what-is-encoding>
- The character stream does conversion from java char to byte representation and vice-versa (as per char encoding).
- The abstract base classes for the character streams are the Reader and Writer class.
- Writer class -- write operation
 - void close() -- close the stream
 - void flush() -- writes data (in memory) to underlying stream/device.
 - void write(char[] b) -- writes char array to underlying stream/device.
 - void write(int b) -- writes a char to underlying stream/device.
- Writer Sub-classes
 - FileWriter, OutputStreamWriter, PrintWriter, BufferedWriter, etc.
- Reader class -- read operation
 - void close() -- close the stream
 - int read(char[] b) -- reads char array from underlying stream/device
 - int read() -- reads a char from the underlying device/stream. Returns -1
- Reader Sub-classes
 - FileReader, InputStreamReader, BufferedReader, etc.

Java NIO

- Java NIO (New IO) is an alternative IO API for Java.
- Java NIO offers a different IO programming model than the traditional IO APIs.
- Since Java 7.
- Java NIO enables you to do non-blocking (not fully) IO.
- Java NIO consist of the following core components:
 - Channels e.g. FileChannel, ...
 - Buffers e.g. ByteBuffer, ...
 - Selectors
- Java NIO also provides "helper" classes Paths & Files.
 - exists()
 - ...

Paths and Files

- A Java Path instance represents a path in the file system. A path can point to either a file or a directory. A path can be absolute or relative.

```
Path path = Paths.get("c:\\data\\myfile.txt");
```

- Files class (Files) provides several static methods for manipulating files in the file system.

```
static InputStream newInputStream(Path, OpenOption...) throws IOException;
static OutputStream newOutputStream(Path, OpenOption...) throws IOException;
static DirectoryStream<Path> newDirectoryStream(Path) throws IOException;
static Path createFile(Path, attribute.FileAttribute<?>...) throws
IOException;
static Path createDirectory(Path, attribute.FileAttribute<?>...) throws
IOException;
static void delete(Path) throws IOException;
static boolean deleteIfExists(Path) throws IOException;
static Path copy(Path, Path, CopyOption...) throws IOException;
static Path move(Path, Path, CopyOption...) throws IOException;
static boolean isSameFile(Path, Path) throws IOException;
static boolean isHidden(Path) throws IOException;
static boolean isDirectory(Path, LinkOption...);
static boolean isRegularFile(Path, LinkOption...);
static long size(Path) throws IOException;
static boolean exists(Path, LinkOption...);
static boolean isReadable(Path);
static boolean isWritable(Path);
static boolean isExecutable(Path);
static List<String> readAllLines(Path) throws IOException;
static Stream<String> lines(Path) throws IOException;
```

Channels and Buffers

- All IO in NIO starts with a Channel.
- A Channel is similar to IO stream.
- From the Channel data can be read into a Buffer.
- Data can also be written from a Buffer into a Channel.

NIO Channels

- Java NIO Channels are similar to IO streams with a few differences:
 - You can both read and write to a Channels. Streams are typically one-way (read or write).
 - Channels can be read and written asynchronously (non-blocking).
 - Channels always read to, or write from, a Buffer.
- Channel Examples
 - FileChannel
 - DatagramChannel // UDP protocol
 - SocketChannel, ServerSocketChannel // TCP protocol

NIO Buffers

- A buffer is essentially a block of memory into which you can write data, which you can then later read again. This memory block is wrapped in a NIO Buffer object, which provides a set of methods that makes it easier to work with the memory block.
- Using a Buffer to read and write data typically follows this 4-step process:
 - Write data into the Buffer
 - Call `buffer.flip()`

- Read data out of the Buffer
 - Call `buffer.clear()` or `buffer.compact()`
- Buffer Examples
 - ByteBuffer
 - CharBuffer
 - DoubleBuffer
 - FloatBuffer
 - IntBuffer
 - LongBuffer
 - ShortBuffer

Channel and Buffer Example

```
RandomAccessFile aFile = new RandomAccessFile("somefile.txt", "rw");
FileChannel inChannel = aFile.getChannel();

ByteBuffer buf = ByteBuffer.allocate(32);

int bytesRead = inChannel.read(buf); // write data into buffer (from channel)
while (bytesRead != -1) {
    System.out.println("Read " + bytesRead);
    buf.flip(); // switch buffer from write mode to read mode

    while(buf.hasRemaining()){
        System.out.print((char) buf.get()); // read data from the buffer
    }

    buf.clear(); // clear the buffer
    bytesRead = inChannel.read(buf);
}
aFile.close();
```

Java NIO vs Java IO

- IO: Stream-oriented
- NIO: Buffer-oriented
- IO: Blocking IO
- NIO: Non-blocking IO

Java 8 Streams

- Java 8 Stream is NOT IO streams.
- `java.util.stream` package.
- Streams follow functional programming model in Java 8.
- The functional programming is based on functional interface (SAM).
- Number of predefined functional interfaces added in Java 8. e.g. Consumer, Supplier, Function, Predicate, ...
- Lambda expression is short-hand way of implementing SAM -- arg types & return type are inferred.

- Java streams represents pipeline of operations through which data is processed.
- Stream operations are of two types

1. Intermediate operations: Yields another stream.

- intermediate operations are again classified as
 1. stateless operation
 - filter(), map(), flatMap(), limit(), skip()
 2. stateful operation
 - sorted(), distinct()

2. Terminal operations: Yields some result.

- reduce()
- forEach()for (Employee e : arr) System.out.println(e);
- collect(), toArray()
- count(), max(), min()
- Stream operations are higher order functions (take functional interfaces as arg).

Java stream characteristics

1. No storage: Stream is an abstraction. Stream doesn't store the data elements. They are stored in source collection or produced at runtime.
2. Immutable: Any operation doesn't change the stream itself. The operations produce new stream of results.
3. Lazy evaluation: Stream is evaluated only if they have terminal operation. If terminal operation is not given, stream is not processed.
4. Not reusable: Streams processed once (terminal operation) cannot be processed again.